Aligned

Dating without guessing games

Berlin-first pitch and technical report for friends, builders, and early advisors.

This document combines the product pitch with a practical technical appendix: what the project needs, what is still missing, what can be copied from existing tools and best practices, and what should be deliberately avoided.

1. Product Pitch

Aligned is a Berlin-first dating platform built for people with ADHD, Autism, and AuDHD who want clearer, calmer, more honest connection. Most dating apps optimise for volume, ambiguity, and superficial attention. Aligned is built around the opposite: structured self-expression, transparent compatibility, lower social guesswork, and fewer but better matches.

The product is not trying to be therapy, diagnosis, or pseudo-psychology. It is a clarity-first dating product for people who are tired of decoding vague profiles, mixed signals, and unclear intent.

Why the problem is real

- Mainstream apps reward performance, ambiguity, and swipe behaviour rather than understanding.
- For neurodivergent users, that creates more emotional fatigue, more hidden incompatibilities, and worse real-world outcomes.
- Berlin is especially relevant because it has strong neurodivergent, creative, and alternative communities, plus a culture that is more tolerant of directness and less tolerant of performative dating behaviour.

How the product works

- Profiles are built around communication style, social energy, sensory preferences, routine versus spontaneity, relationship intent, and lifestyle rhythm.
- Matches are explained in plain language: likely strengths, likely friction points, and confidence level.
- Conversation support reduces first-message friction without pretending to predict chemistry.

Why NLP matters

The NLP layer is used for practical, bounded tasks that improve outcomes. It helps turn messy free text into usable signals, highlights low-signal or contradictory profiles, improves match explanations, and supports moderation. It is not used to diagnose users, infer mental health severity, or claim deep psychological certainty from limited data.

Why this niche can win

- High pain, high intent, weak current alternatives.
- A credible paid model already exists in dating; the opportunity is to improve outcomes, not invent monetisation from scratch.
- The niche is focused enough to differentiate, but large enough to scale city by city if density works.

How it scales

- Phase 1: win Berlin inside the Ring by creating meaningful density and repeatable match quality.

- Phase 2: expand to similar cities with compatible user density and culture.
- Phase 3: compound through referrals, creator/community partnerships, and outcome-led reputation.

2. Technical Report: What the Project Still Needs

This section is intentionally practical. It focuses on the things the project does not currently have but should add if the goal is fast execution, legal defensibility, better trust, and higher probability of meaningful success.

Area	**Current assumption**	**Missing / should be added**
Core app	Concept and product direction are defined.	Production-grade web app, API, database schema, auth, admin
AI/NLP	Idea is clear.	Bounded prompt layer, schema-based outputs, eval set, modera
Compliance	Awareness of GDPR/DSA/AI boundaries exists.	Privacy policy, consent design, records of processing, data rete
Growth	Berlin-first logic exists.	Waitlist system, referral loop, CRM, creator/community outreach
Trust & safety	Policy direction exists.	Moderation queue, report handling, incident logging, banned-pa

Recommended stack to copy rather than invent

- **Frontend:** Next.js or React + TypeScript, Tailwind CSS, React Hook Form, Zod, TanStack Query.
- **Backend:** FastAPI or Node/TypeScript API. Pick one, not both. FastAPI is strong for typed service layers and moderation/NLP pipelines; Node is strong if you want one-language full stack.
- **Database:** PostgreSQL with pgvector only if vector search is genuinely needed. Avoid a separate vector database early.
- **Auth and data platform:** Supabase is the most practical all-in-one starting point if speed matters.
- **Analytics and experiments:** PostHog for product analytics, session replay, feature flags, and experiments.
- **Error monitoring:** Sentry from day one.
- **LLM layer:** OpenAI Responses API with structured outputs and explicit schemas for any user-facing NLP.

Libraries and tools worth using

Category	Recommended tools	Why they help
Forms + validation	React Hook Form, Zod	Fast onboarding, strict input validation, fewer broken payloads.
State + fetching	TanStack Query	Cleaner server-state handling and caching.
UI	Tailwind CSS, shadcn/ui or Radix	Fast, consistent UI without design chaos.
Database	PostgreSQL, pgvector	Single source of truth; vector search only if later justified.
Auth	Supabase Auth or Clerk	Off-the-shelf auth, sessions, email magic links, less security foo
Analytics	PostHog	Funnels, retention, session replay, feature flags, experiments.
Monitoring	Sentry	Errors, performance traces, release tracking.
Search	Postgres full-text first; Meilisearch later if needed	Do not overbuild search in v1.
Queue/jobs	Celery/RQ (Python) or BullMQ (Node)	Async moderation, summaries, notifications later.
Email/CRM	Resend / Loops / Customer.io	Waitlist, invites, lifecycle messaging.
Payments	Stripe later, not day one	Delay until product value is obvious.

NLP system design: what to build, what not to build

- **Build:** profile summarisation, communication-style extraction, low-signal detection, contradiction checks, safer first-message suggestions, moderation assistance.
- **Do not build early:** personality scoring, mental-health severity inference, diagnosis verification, manipulative engagement optimisation, deep recommendation models with fake precision.
- **Use schemas:** every model output should land in strict JSON or typed structures. This is where structured outputs matter most.
- **Have confidence levels:** high, medium, low. When the signal is weak, the system should say so.

Training, data, and evaluation requirements

- **No model training is required early.** The project can move fast with API-based LLMs, structured prompts, and a strong evaluation set.
- **Create a labelled eval set.** At least 100–200 synthetic and manually reviewed profiles/chats covering high clarity, low clarity, contradictions, fetishising language, coercive tone, and harmless edge cases.
- **Human review is non-negotiable.** At least one person should manually review outputs before trusting any moderation or compatibility signal.
- **Prompt/version control is required.** Store prompts, schemas, evaluations, and example failures in the repo.
- **Research usage should stay at product-design level.** Use validated evidence to inform question design and safety boundaries, not to overclaim scientific matching.

Compliance, trust, and safety requirements

Because the app touches health-adjacent data and runs as an online platform, compliance is not optional. It is part of product quality.

What must be added early

- Explicit consent design for sensitive profile data.
- A privacy policy and clear data-retention policy.
- Delete/export capability in the roadmap from the start.
- A reporting and appeals flow for harmful interactions.
- Incident logging and internal moderation notes.
- A written AI boundaries policy explaining what the system does and does not infer.

Security and operations

- Passwordless or hardened auth flow, rate limiting, abuse controls, and audit logging.
- Environment variable hygiene, secret management, and least-privilege database roles.
- Backups, staging environment, and release rollback plan.
- CI checks for type safety, tests, linting, and prompt-schema validation.

People and process requirements that do not exist yet

- **Legal reviewer:** GDPR/privacy/terms review before public beta.
- **Trust & safety owner:** even if part-time, someone must own reports and moderation policy.
- **Design reviewer:** calm UX matters here; a chaotic interface will hurt adoption.
- **Early user researcher:** 10–20 structured interviews in Berlin before and during beta.
- **Ops discipline:** weekly cohort review, match quality review, and failure log.

Recommended 90-day technical sequence

Weeks 1–2	Repo, product docs, auth, DB schema, onboarding, profile system, waitlist, analytics wiring.
Weeks 3–4	Matching v1, chat v1, profile summaries, low-signal nudges, admin seed tools.
Weeks 5–6	Moderation queue, report flow, confidence labels, session replay, error tracking, legal draft revie
Weeks 7–8	Closed Berlin alpha, manual seeding, user interviews, onboarding and ranking fixes.
Weeks 9–12	Berlin beta, measured iteration, first premium hypotheses, referral loop tests.

What should be deliberately postponed

- Native mobile app.
- Heavy custom model training or GPU workflows.
- Separate vector infrastructure unless retrieval genuinely becomes core.
- Complicated social/community features.
- Aggressive monetisation before user outcomes are proven.

- Anything that turns the product into therapy, diagnosis, or pseudo-science.

Fast budget view for missing pieces

Category	Lean estimate	Comment
Legal/compliance	€4k–€6k	Privacy, terms, consent flow, platform obligations.
Infra + tools	€1k–€2k	Hosting, DB, analytics, monitoring, email.
LLM/API usage	€500–€1.5k	Enough for MVP and early beta if prompts are disciplined.
Design and UX help	€1k–€3k	Optional but high leverage if interface quality is weak.
User research / acquisition	€2k–€4k	Berlin seeding, creators, community outreach, events.

Bottom line

The project does not need custom AI research, heavy training infrastructure, or a huge team to get started. It needs a disciplined product stack, clear trust boundaries, proper analytics, a compliance baseline, and strong operational discipline in Berlin. The fastest path is to copy proven infrastructure where it helps, reserve AI for bounded, high-value tasks, and avoid building anything that creates false certainty.

If Aligned becomes strong at clarity, trust, and real-world outcomes before it tries to become sophisticated, it has a real chance.

Source notes

Technical recommendations in this report draw on current official documentation for Supabase, PostHog, Sentry, and OpenAI, plus current EU platform and data-protection rules. Legal review is still required before any public launch.